

Mikhail Yurochkin

✉ moonfolk@umich.edu • 🌐 <https://moonfolk.github.io/>

Research Interests

- Algorithmic Fairness
- Model Fusion and Federated Learning
- Applications of Optimal Transport in Machine Learning
- Bayesian (nonparametric) modeling and inference

Education

Ph.D., Statistics 2013 – 2018

University of Michigan, Ann Arbor, MI, United States

- Thesis: “Geometric Inference in Bayesian Hierarchical Models with Applications to Topic Modeling”
- Advisor: Professor XuanLong Nguyen

M.A., Statistics 2013 – 2015

University of Michigan, Ann Arbor, MI, United States

Bachelor's, Applied Mathematics and Physics 2009 – 2013

Moscow Institute of Physics and Technology, Moscow, Russia

Work Experience

Research Staff Member June 2018 – Present

IBM Research and MIT-IBM Watson AI Lab, Cambridge, MA, United States

Machine Learning and Artificial Intelligence research.

- Leading a team of researchers and engineers on a Federated Learning project.
- Leading a Tech for Justice project on developing software and demonstrations for Individual Fairness algorithms developed in collaboration with Yuekai Sun's group at the University of Michigan.
- Leading a project with Justin Solomon's Geometric Data Processing group at MIT on new applications of Optimal Transport to machine learning problems.
- Collaborating with Tamara Broderick and her group at MIT on Gaussian Processes for Big Data.

Research Associate May 2018 – June 2018

University of Michigan, Ann Arbor, MI, United States

Leading two research projects in collaboration with senior PhD students and faculty. This work led to two paper submissions to NeurIPS 2018. Contributing to project on analyzing driving behaviors using Bayesian modeling as part of collaboration between University of Michigan and Toyota Research Institute.

Data science research intern June 2017 – Aug 2017

Adobe, San Jose, CA, United States

Explored various approaches for graph learning in the context of convolutional neural networks on graphs. Developed novel approach for deep learning on graph structured data which can automatically learn the latent graph representation. Applied developed techniques to forecasting number of visits to a major retailer's website across cities for improved marketing and sales. Submitted a patent application based on these results.

Consultant for Statistics, Computing, and Analytics Research Sept 2016 – Dec 2017

University of Michigan, Ann Arbor, MI, United States

Individual appointments and walk-in consultations for faculty and graduate students. Assisted researchers in areas such as biology, dentistry, marketing, survey methodology, political science, public health, computer science, and more with identifying appropriate statistical methodology for the corresponding data and research question. Additionally helped with data collection and processing, software choice and implementation.

Science team intern

May 2016 – Aug 2016

LogicBlox/Predictix, Atlanta, GA, United States

Developed novel machine learning approaches for analyzing sales data of a major retailer. Applied expertise in Bayesian modeling to propose, implement and test explainable retail demand forecasting model based on Factorization Machines and Indian Buffet Process. Published a paper at NeurIPS 2017 based on this work.

Graduate Student Research Assistant

July 2014 – April 2018

University of Michigan, Ann Arbor, MI, United States

Conducting research for the NSF-supported projects:

- “Geometric approaches to hierarchical and nonparametric model-based inference” (award #1351362). Formulated topic modeling with Latent Dirichlet Allocation as geometric problem. Developed two new geometry based algorithms for topic modeling, significantly outperforming previous state-of-the-art techniques in terms of estimation speed. Second algorithm additionally can learn latent number of topics utilizing geometric structure of the problem. These results were published at NeurIPS 2016 and 2017. Gave multiple talks and poster presentations. My adviser included some of these results in the curriculum of a data science masters level class and a research oriented undergraduate level class.
- “Data is Social: Exploiting Data Relationships to Detect Insider Attacks” (award #1409551). Proposed Bayesian modeling approaches to identify insider attacks from SQL queries of a bank database. Worked in collaboration with PhD students and faculty from the University at Buffalo.

Graduate Student Instructor

Sept 2013 – April 2016

University of Michigan, Ann Arbor, MI, United States

Teaching labs, grading, holding office hours:

- Topics in Biostatistics Jan 2015 – April 2016
- Applied Probability Sept 2014 – Dec 2014
- Introduction to Statistics Jan 2014 – April 2014
- Introduction to Probability Sept 2013 – Dec 2013

Publications

- [1] F. Petersen, D. Mukherjee, Y. Sun, and **M. Yurochkin**. Post-processing for Individual Fairness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [2] M. Agarwal, **M. Yurochkin**, and Y. Sun. On sensitivity of meta-learning to support data. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [3] S. Maity, D. Mukherjee, **M. Yurochkin**, and Y. Sun. Does enforcing fairness mitigate biases caused by subpopulation shift? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [4] V. Huynh, N. Ho, N. Dam, X. Nguyen, **M. Yurochkin**, H. Bui, and D. Phung. On efficient multilevel Clustering via Wasserstein distances. *Journal of Machine Learning Research (JMLR)*, 2021.
- [5] D. Mukherjee, A. Guha, J. Solomon, Y. Sun, and **M. Yurochkin**. Outlier-Robust Optimal Transport. In *International Conference on Machine Learning (ICML)*, 2021.
- [6] S. Maity, S. Xue, **M. Yurochkin**, and Y. Sun. Statistical inference for individual fairness. In *International Conference on Learning Representations (ICLR)*, 2021.
- [7] A. Bower, H. Eftekhari, **M. Yurochkin**, and Y. Sun. Individually Fair Rankings. In *International Conference on Learning Representations (ICLR)*, 2021.
- [8] A. Vargo, F. Zhang, **M. Yurochkin**, and Y. Sun. Individually Fair Gradient Boosting. In *International Conference on Learning Representations (ICLR)*, 2021. Spotlight.
- [9] **M. Yurochkin** and Y. Sun. SenSel: Sensitive Set Invariance for Enforcing Individual Fairness. In *International Conference on Learning Representations (ICLR)*, 2021. Oral.
- [10] M. Weber, **M. Yurochkin**, S. Botros, and V. Markov. Black Loans Matter: Distributionally Robust Fairness for Fighting Subgroup Discrimination. *NeurIPS Fair AI in Finance Workshop*, 2020. Spotlight Talk.

- [11] L. Li, A. Genevay, **M. Yurochkin**, and J. Solomon. Continuous Regularized Wasserstein Barycenters. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [12] S. Clatici*, **M. Yurochkin***, S. Ghosh, and J. Solomon. Model Fusion with Kullback–Leibler Divergence. In *International Conference on Machine Learning (ICML)*, 2020.
- [13] D. Mukherjee*, **M. Yurochkin***, M. Banerjee, and Y. Sun. Two Simple Ways to Learn Individual Fairness Metric from Data. In *International Conference on Machine Learning (ICML)*, 2020.
- [14] S. Xue, **M. Yurochkin**, and Y. Sun. Auditing ML models for individual bias and unfairness. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2020.
- [15] H. Wang, **M. Yurochkin**, Y. Sun, D. Papailiopoulos, and Y. Khazaeni. Federated Learning with Matched Averaging. In *International Conference on Learning Representations (ICLR)*, 2020. Oral.
- [16] **M. Yurochkin**, A. Bower, and Y. Sun. Training individually fair ML models with sensitive subspace robustness. In *International Conference on Learning Representations (ICLR)*, 2020. Spotlight.
- [17] **M. Yurochkin**, S. Clatici, E. Chien, F. Mirzazadeh, and J. Solomon. Hierarchical Optimal Transport for Document Representation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [18] **M. Yurochkin**, M. Agarwal, S. Ghosh, K. Greenewald, and N. Hoang. Statistical Model Aggregation via Parameter Matching. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [19] P. Monteleiller, S. Clatici, E. Chien, F. Mirzazadeh, J. Solomon, and **M. Yurochkin**. Alleviating Label Switching with Optimal Transport. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [20] **M. Yurochkin**, Z. Fan, A. Guha, P. Koutris, and X. Nguyen. Scalable inference of topic evolution via models for latent geometric structures. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [21] **M. Yurochkin**, A. Guha, Y. Sun, and X. Nguyen. Dirichlet Simplex Nest and Geometric Inference. In *International Conference on Machine Learning (ICML)*, 2019. Long Talk.
- [22] **M. Yurochkin**, M. Agarwal, S. Ghosh, K. Greenewald, N. Hoang, and Y. Khazaeni. Bayesian Nonparametric Federated Learning of Neural Networks. In *International Conference on Machine Learning (ICML)*, 2019.
- [23] **M. Yurochkin**, S. Upadhyay, D. Bouneffouf, M. Agarwal, and Y. Khazaeni. Online Semi-Supervised Learning with Bandit Feedback. *ICLR Limited Labeled Data (LLD) Workshop*, 2019.
- [24] **Mikhail Yurochkin**. *Geometric Inference in Bayesian Hierarchical Models with Applications to Topic Modeling*. PhD thesis, University of Michigan, 2018.
- [25] **M. Yurochkin**, A. Guha, and X. Nguyen. Conic Scan-and-Cover algorithms for nonparametric topic modeling. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [26] **M. Yurochkin**, X. Nguyen, and N. Vasiloglou. Multi-way Interacting Regression via Factorization Machines. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [27] N. Ho, X. Nguyen, **M. Yurochkin**, H. Bui, V. Huynh, and D. Phung. Multilevel Clustering via Wasserstein Means. In *International Conference on Machine Learning (ICML)*, 2017.
- [28] **M. Yurochkin** and X. Nguyen. Geometric Dirichlet Means algorithm for topic inference. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2016.

Patents

- **M. Yurochkin**, D. Mukherjee, M. Banerjee, Y. Sun, and S. Upadhyay. Learning Mahalanobis Distance Metrics from Data. Filed on June 11, 2021.
- S. Upadhyay, **M. Yurochkin**, D. Mukherjee, Y. Sun, A. Bower, H. Eftekhari, A. Vargo, and F. Zhang. Training Individually Fair Machine Learning Algorithms via Distributionally Robust Optimization. Filed on March 25, 2021.
- Y. Khazaeni, E. Daly, C. Muise, and **M. Yurochkin**. Artificial Intelligence for Learning Path Recommendations. Filed on December 31, 2020.
- K. Greenewald, **M. Yurochkin**, M. Agarwal, S. Ghosh, N. Hoang and Y. Khazaeni. A method for combining pre-trained neural networks into a memory and computation efficient global model. Filed on September 20, 2019.
- S. Upadhyay, **M. Yurochkin**, M. Agarwal, D. Bouneffouf and Y. Khazaeni. Method for Online Partially Rewarded Learning. Filed on August 28, 2019.

Talks

Invited Talks.....

- Black Loans Matter: Fighting Bias for AI Fairness in Lending
@ KDD Workshop on Machine Learning in Finance (Keynote) 2021
- Model fusion via single-round FL
@ Enterprise-Strength Federated Learning, ICML Expo 2021
- Practical Individual Fairness algorithms
@ ODSC West 2021
@ Foundations of Algorithmic Fairness (FAF) workshop 2021
@ Research Talks for Enel 2021
@ IBM Research, Cambridge Lab All-hands 2021
- Fusion of Heterogeneous Models in Federated Learning
@ The 2nd Annual Federated & Distributed/ Decentralized Machine Learning Conference 2021
- Federated Learning: Practice and Modern Algorithms
@ ODSC Europe 2021
- Invited panelist
@ Federated Learning, AI and Data Security Summit, HUB Security event 2021
- Fairness in mortgage lending
@ Tech for Racial and Social Justice, UW-Madison Data Science Bazaar 2021
- Individual Fairness through Robustness
@ Algorithmic fairness with statistical guarantees, CMStatistics session 2020
@ Global Data Science Community meet-up, UniCredit event 2020
- Fairness in AI
@ What's Next in AI: AI we can trust, MIT-IBM conference 2020
- Bayesian nonparametric fusion of heterogeneous models
@ IBM Research Cambridge seminar series 2019
- Geometric Inference in Bayesian Hierarchical Models with Applications to Topic Modeling
@ MIT Geometric Data Processing Group seminar 2018

Conference Presentations.....

- Data Science Research Forum (poster) 2017
- Joint Statistical Meetings (speed session talk and poster) 2017
- Michigan Student Symposium for Interdisciplinary Statistical Sciences (talk) 2017
- Michigan Institute for Computational Discovery and Engineering Symposium (poster) 2016
- Michigan Student Symposium for Interdisciplinary Statistical Sciences (poster) 2016
- From Industrial Statistics to Data Science (poster) 2015

Awards

- Outstanding Technical Achievement Award (IBM internal) 2021
- EB1 Green Card (outstanding professor or researcher) 2021
- A-level Research Accomplishment (IBM internal) 2020
- Faces of IBM: AI Pioneer (IBM internal) 2020
- Outstanding Demonstration - Honorable Mention award NeurIPS 2020
- Reviewer award NeurIPS 2018, 2019
- Rackham Conference Travel Grant 2016, 2017
- Abramov scholarship (awarded to less than 5% students at MIPT) 2010

Mentoring

- Luann Jung, MIT Masters of Engineering (MEng) student (with J. Solomon) 2021-present
- Yining Chen, MIT-IBM summer intern (with J. Solomon and Y. Sun) 2021
- Haingoharijao Faniriniaina Ramandiamanana, African Master in Machine Intelligence (AMMI) student (with J. Solomon and D. Palmer) 2020
- Debarghya Mukherjee, MIT-IBM summer intern (with J. Solomon and Y. Sun) 2020
- Hongyi Wang, IBM Research summer intern 2019

Professional Activities

Organization and Service.....

- Practical Bayesian Methods for Big Data and Big Models workshop (co-organizer) @ MIT-IBM AI week 2019
- Optimal Transport seminar (with Y. Mroueh) @ IBM Research AI seminar series 2019

Reviewer.....

- Neural Information Processing Systems (NeurIPS) 2016-2021
- International Conference on Learning Representations (ICLR) 2018-2020
- International Conference on Machine Learning (ICML) 2017-2019
- Bayesian Analysis (BA) 2019
- Journal of Machine Learning Research (JMLR) 2019
- Journal of Computational and Graphical Statistics (JCGS) 2016

Media coverage

- "New research helps make AI fairer in decision-making", IBM Research 2021
- "Finding a good read among billions of choices", MIT News 2019
- "Optimal Transport for Label Switching: Using Geometry to Solve Problems in AI", IBM Research 2019

Additional information

- Personal webpage: <https://moonfolk.github.io/>
- LinkedIn: <https://www.linkedin.com/in/mikhail-yurochkin-a45659114>
- GitHub: <https://github.com/moonfolk>
- Google Scholar: <https://scholar.google.com/citations?user=QjBF9sUAAAAJ>